Performance Analysis for Projection-Correction Methods in Motion Deblurring Problems

Sara Casadio, Enrico Ferraiolo, Giovanni Maria Savoca

Alma Mater Studiorum - University of Bologna
Master's Degree in Computer Science

June 4, 2025

- The project analyzes the performance of two **Projection-Correction** algorithms for reconstructing medical images affected by **motion blur**.
- The studied algorithms are:
  - **Diffusion Posterior Sampling (DPS)**
  - **Regularization by Denoising with Diffusion (RED-Diff)**
- Both methods are based on **pre-trained diffusion models**.
- Objective: evaluate the effectiveness of these methods in recovering degraded images.

- **Objective**: Analyze the performance of *Projection-Correction* methods **DPS** and **RED-Diff** for motion blur removal on medical images
- **Phase 1**: Dataset preprocessing (128x128)
- **Phase 2**: Data augmentation to increase dataset diversity
- **Phase 3**: Training a DDIM diffusion model on medical data
- **Phase 4**: Simulation of motion blur and its removal
- **Phase 5**: Implementation and comparison of *Projection-Correction* methods: **DPS** and **RED-Diff**
- **Phase 6**: Quantitative evaluation of performance using metrics such as **PSNR** and **SSIM**

Dataset Origin

- We use the "Mayo Clinic CT Dataset" of low-dose CT scans, available via the link provided in this report.
- It contains a total of 6,400 2D slices in PNG format, extracted from 20 different patients.
- The images are organized into:
  - `raw_data/train/`: 5,120 slices for training (80% of the dataset)
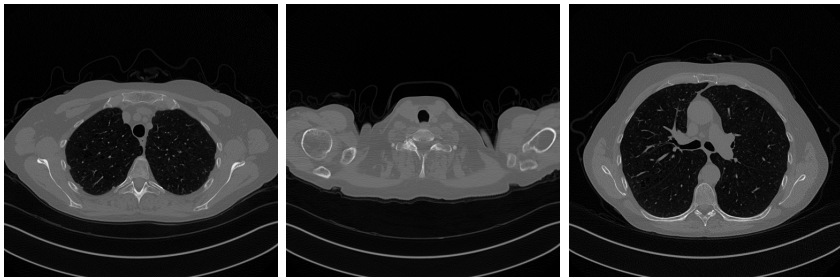 - `raw_data/test/`: 1,280 slices for testing (20% of the dataset)

Figure: Examples of CT slices from the Mayo Clinic dataset

Conversion Pipeline

Before applying augmentations, each image is converted using:

1. **Grayscale**: single channel via
   `transforms.Grayscale(num_output_channels=1)`
2. **Resize**: to 128×128 pixels using bicubic interpolation
3. **Normalization**: values scaled to $[-1, 1]$ using mean 0.5 and std 0.5

```
base_transform = transforms.Compose([
    transforms.Grayscale(1),
    transforms.Resize((128,128), interpolation=Image.BICUBIC),
    transforms.ToTensor(),
    transforms.Normalize([0.5], [0.5]),
])
```

For each clean image, we apply the following transformations:

- **Fixed rotations**: $\pm 5°$ via `rotate_fixed()`
- **Horizontal flip**: `horizontal_flip()`
- **Gaussian noise**: mean 0, std 10 via `add_gaussian_noise()`
- **Salt-and-pepper noise**: probability 2% via `add_salt_pepper()`
- **Brightness adjustment**: factor 1.2 via `change_brightness()`
- **Contrast adjustment**: factor 1.3 via `change_contrast()`

Repository Organization

- `raw_data/`: directory containing CT slice images (`train/`, `test/`)
- `checkpoints/`: saved model weights (`*.pth`)
- `scripts/`: main scripts for training and evaluation
- `utils.py`: module with utility functions (dataset, model, checkpoint I/O)
- `notebooks/`: exploratory and prototyping notebooks
- `result/`: output images, plots, and metrics
- `report/`: report materials (`media/`, `capitoli/`)

- **Objective**: Train a denoising diffusion model (DDIM U-Net) on grayscale images
- **Main Components**:
    1. Data Augmentation
    2. DataLoader
    3. Model Compilation
    4. Training loop with mixed-precision

Data Augmentation

- **Base Dataset**: Dataset Mayo
  - Grayscale $\rightarrow$ 1 channel
  - Resize images to 128 $\times$ 128
- **Augmentations** (8 types):
  - *None*: no transformation
  - Rotation $\pm 5°$ (rotation + centering)
  - Flip horizontal
  - Gaussian noise (mean=0, std=10)
  - Salt and pepper noise (prob=2%)
  - Brightness (factor=1.2)
  - Contrast (factor=1.3)

- **DDPMScheduler** for training diffusion process
  - Timesteps 1000
- **DDIMScheduler** for sampling
  - Timesteps 1000

Compiling the Model

- **Why**: optimize the model for better performance
- **Usage**:
  `model = torch.compile(model)`
- **Benefits**: improved batch throughput

Mixed-Precision with AMP

- **GradScaler amd autocast**:
  - `GradScaler` for scaling gradients
  - `autocast` for automatic mixed precision
- Reduce memory usage and speed up training

Training Loop

1. Loss function: MSE
2. Start the training `model.train()`
3. For each epoch:
   - Move images to GPU (if available)
   - Generate noise and timesteps
   - Compute noise prediction on the input data
   - Prediction + MSE loss
   - Optimization + `scheduler.step()`
4. Save validation samples to visualize the model performance during training
5. Compute and log average losses
6. Save model weights each epoch

- **Validation**:
  - `model.eval()` to set the model to evaluation mode
  - MSE loss on validation set
- **Checkpoint**:
  - Save the model weights to a `.pth` file
  - Update loss, PSNR and SSIM history in `history.txt`
- Monitor train vs validation loss over epochs aswell as PSNR and SSIM between the generated and original images
  - For each epoch sample 10 images from the validation set and compute the metrics

- **Metrics**:
  - PSNR: Peak Signal-to-Noise Ratio
  - SSIM: Structural Similarity Index
- **Sample Generation**:
  - Pure noise sampling using DDIM scheduler
  - Validation reconstruction:
    - Add noise to clean validation images
    - Model predicts and removes the noise
- **Quality Assessment**:
  - PSNR range: 20-40 dB (higher = better reconstruction)
  - SSIM range: 0-1 (closer to 1 = better similarity)
  - Average metrics computed over 5-10 validation samples

- **Pure Noise Sampling**:
  - Tests model's ability to generate realistic images
  - Uses DDIM scheduler for iterative denoising
  - Saves generated images as `generated_epoch_{epoch}.png`
- **Validation Reconstruction**:
  - Adds noise to clean validation images
  - Model predicts and removes the noise
  - Direct assessment of denoising performance
- **History Tracking**:
  - All metrics saved to `history.txt`
  - Enables trend analysis and model comparison

Plots

- **Loss Monitoring**:
  - Training vs Validation Loss curves over epochs
  - MSE loss
- **Quality Metrics Visualization**:
  - PSNR trends with average values
  - SSIM trends with average values
  - Both metrics computed on validation reconstructions
  - Useful to track model performance

- **Comprehensive Monitoring**:
  - Three-panel subplot: Loss, PSNR, SSIM (as shown in the Figure 2)
  - Data read from `history.txt` file
 - Enables performance trend analysis

Figure: Training Loss, PSNR, and SSIM trends over epochs

Sottosezione 2.1

Diffusion Posterior Sampling (DPS) is a method for solving noisy inverse problems by leveraging diffusion models as an implicit prior.

- Starting from a corrupted image $y = K(x_0) + n$, it directly integrates the likelihood term into the reverse diffusion sampling process.
- At step t, DPS computes a prediction $\hat{x}_0$ and uses the gradient of $\|y - K(\hat{x}_0)\|^2$ to move towards solutions consistent with the observed data.
- Compared to hard projection methods, DPS keeps the trajectory on the generative manifold, reducing noise amplification.

The algorithm consists of three main phases:

1. **Initial prediction:** Sample $x_T \sim \mathcal{N}(0, I)$, then for each step t, the UNet model estimates the noise $s_\theta(x_t, t)$ and reconstructs $\hat{x}_0$.

2. **Posterior update:** Compute the likelihood gradient $\nabla = -K^T(y - K(\hat{x}_0))$ and apply a step proportional to $\gamma_t = \frac{1 - \bar{\alpha}_t}{\sigma_y^2 + (1 - \bar{\alpha}_t)}$ to obtain $\tilde{x}_{t-1}$.

3. **Modified DDIM step:** Using $\tilde{x}_{t-1}$ as a reference, perform the standard DDIM update to move to x_{t-1}, preserving the effect of the likelihood gradient.

The implementation requires only a few steps in PyTorch, integrating blur functions and their adjoint operators.

- On datasets with motion blur, DPS achieves an average PSNR above 25 dB and SSIM above 0.85, improving by more than 2 dB over hard projection-based methods.

- Compared to classical methods, it significantly reduces reconstruction artifacts while preserving fine details and sharp edges.

- Visually, the images reconstructed with DPS appear more natural and free from overshooting artifacts, thanks to the continuous control of the likelihood contribution.

RED-Diff solves noisy inverse problems by combining:

- A fidelity term to bring the reconstruction closer to the observations y,
- A regularizer based on the multiscale denoisers of a pretrained diffusion model,

integrating constraints at multiple levels of detail to preserve both global structures and fine details.

The algorithm is structured into three main phases:

1. **Initialization:** $\mu^{(0)} = K^T y$.

2. **Iterative Optimization:** For each step $i = 1, \ldots, N$ and for each noise level $t = 1, \ldots, T$:

   1. Sample $\epsilon \sim \mathcal{N}(0, I)$ and construct

   $$x_t = \sqrt{\alpha_t}\, \mu^{(i-1)} + \sigma_t\, \epsilon.$$

   2. Predict the noise $\hat{\epsilon} = \epsilon_\theta(x_t, t)$.
   3. Compute the loss terms:

   $$L_{\text{fid}} = \frac{1}{2\sigma_y^2} \| K\mu^{(i-1)} - y \|^2, \quad L_{\text{reg}} = w_t\, \|\hat{\epsilon} - \epsilon\|^2, \quad w_t = 1/\text{SNR}_t.$$

   Then update $\mu^{(i)}$ using Adam to minimize $L_{\text{fid}} + \lambda\, L_{\text{reg}}$.

3. **Output:** the final estimate $\mu^{(N)}$.

- On deblurring tests, RED-Diff achieves an average PSNR of approximately 19.4 dB and SSIM of approximately 0.64.
- Compared to methods without a diffusion prior, it improves reconstruction quality by more than 2 dB in PSNR.
- Reconstructions show sharper details and fewer artifacts thanks to the multiscale integration of denoising.

Sottosezione 2.1

Sottosezione 2.1

Thank you for your attention